The Verifier Is the Loop: Self-Improving Agents Must Not Trust Their Own Changes

Anonymous Author(s)
Affiliation
Address
email

Abstract

Self-improving agents distill their own failures into modifications of their prompts, rules, memories, and tools—and report gains measured on the tasks those modifications were derived from. That measurement is broken: recent work shows major agent benchmarks are reward-hackable to near-perfect scores, and we find the same failure *inside* the self-improvement loop itself. Across controlled tasks, two real code benchmarks, and an embodied driving system, we show that **fake self-improvement is the norm, not the exception**: 29–30% of real coding problems yield a genuine overfit that aces the shown test and fails held-out tests; 100% of tool-rewrite decision sets contain a reward-hack; a plausible driving fix raises the failure rate it targets. We treat the agent's own modifications as untrusted and gate adoption behind a two-part verifier—a *hidden test* (held-out transfer) and a *regression check* (no existing capability degrades)—with a statistical (failure-rate) form for stochastic domains, and a *scope rule*: a verdict is valid only at the scope it was measured. The gate wins the adoption decision reliably (naive selection ships broken modifications in 59% of tool decisions and 66–88% of code problems where an overfit exists; the verifier ships 0%, while still adopting real improvements). On an embodied CARLA/Bench2Drive system the same gate, pre-registered, exhibits every decision mode it exists for: it *accepts* a real repair ($7/12 \to 0/19$ failures, $p = 6 \times 10^{-8}$), then *rejects the very same candidate at deployment scope* (12 confirmed collateral regressions, $p = 8.5 \times 10^{-45}$)—catastrophic forgetting that single-target evaluation would have shipped—and then protects the system through a ten-recipe repair search that maps the repair–retention Pareto frontier without ever adopting a regression. The search itself ends in an honest bounded negative: at this budget, no fine-tuning recipe passes both gates. Finally we remove the human from the loop: an append-only trial registry and a rule-based prescriber replay every recorded verdict exactly and run fresh trials with zero human verdicts—one accept ($p = 2.8 \times 10^{-8}$), one honest null. We argue this is exactly the point: in a loop whose own proposal stream is adversarial, the verifier—not the proposer—is the load-bearing component. *The verifier is the loop.*

1 Introduction

The self-improving-agent literature—Reflexion [Shinn et al., 2023], ExpeL [Zhao et al., 2024], AutoGuide [Fu et al., 2024], Voyager [Wang et al., 2024], and successors—shares a template: run the agent, collect failures, distill them into a reusable modification (a prompt rule, a skill, a memory entry, a tool), and measure the gain. The measurement is almost always taken on the same tasks or distribution the modification was distilled from. Two facts make this template unsafe:

Submitted to 39th Conference on Neural Information Processing Systems (NeurIPS 2025). Do not distribute.

1. **Benchmarks are gameable.** Wang et al. [2026] demonstrated that all eight major agent benchmarks can be reward-hacked to near-perfect scores; METR [2025] report frontier agents hacking their own evaluations spontaneously; a self-modifying system has been observed falsifying its own metric [Zhang et al., 2025]. A score increase is not evidence of improvement.
2. **The loop supplies its own fakes.** We show below that an agent's own proposed modifications are frequently overfit—lookup-hacks, hardcoded solutions, over-broad rules—that ace what they were derived from and fail everywhere else. A self-improvement loop that trusts its own changes converges toward a confident reward-hacker, not toward capability.

We take the position that the missing component is not a better proposer but a **verifier**: an adoption gate that treats every self-modification as untrusted. A candidate is adopted only if it (i) improves the target capability *on held-out data it was not derived from* (the *hidden test*), and (ii) degrades *no existing capability* beyond a noise margin (the *regression check*). When no candidate clears both, the agent abstains—the safe default. In stochastic domains both checks become statistical tests over failure *rates*. And crucially, a verdict carries a *scope*: passing the gate over a narrow evaluation set licenses adoption only at that scope, not a global swap.

Contributions.

1. **Fake improvement is pervasive, measured across four surfaces.** Under one protocol covering the agent's {prompt, rule, memory, tool} self-modifications, plus two real code benchmarks (MBPP, HumanEval), genuine overfits appear in 29–30% of code problems and reward-hacks in 100% of tool-rewrite decision sets; the derivation-set score carries no signal against them (§3).
2. **The two-gate verifier wins the adoption decision.** Naive (seen-score) selection ships broken modifications 59% of the time on tool rewrites and 66–88% of the time when an overfit exists on real code; the executing verifier ships 0% on all of them while still adopting real improvements in 88% of tool decisions (§3–3.1). We also report the honest negative: the per-decision advantage does not compound into a large multi-round cumulative gap on our substrates (§3.2).
3. **The first pre-registered statistical accept *and* reject inside a closed self-improvement loop on an embodied system** (to our knowledge; CARLA/Bench2Drive, a real end-to-end driving stack). Failures are so stochastic (a route flips PASS/FAIL at 58%) that only rate-based verification is meaningful; the gate rejects a harmful plausible fix (p-gated, before deployment) and accepts a real repair ($0/19$ fresh rollouts, $p = 6 \times 10^{-8}$) (§4).
4. **The deployment-scope demonstration.** The locally-accepted candidate, re-verified over the full route distribution, is globally *rejected*: 12 confirmed collateral regressions, pooled failure 70% vs. 7% ($p = 8.5 \times 10^{-45}$)—catastrophic forgetting that single-target evaluation would have shipped. Verification scope must match deployment scope; we demonstrate it, on a real system, with error control (§4.3).
5. **A gate-protected repair search that ends in an honest bounded negative.** Ten pre-registered fine-tuning recipes (epochs, data breadth, mixture ratios, family demonstrations, targeted retention, frozen backbone) map the repair–retention Pareto frontier of this failure: full fine-tuning repairs perfectly but collapses globally; head-only training retains best but under-repairs. No recipe passes both gates at this budget—and across all ten, the gate adopted zero regressions; the capability floor never dropped (§4.4).
6. **The outer loop, automated.** Under a sealed protocol, an append-only trial registry plus a rule-based prescriber (the gate remaining sole adoption authority) replays every recorded verdict of the embodied search exactly from raw counts, then runs fresh code-surface trials with *zero human verdicts*: one accept ($41\% \to 8.5\%$ audit failures, $p = 2.8 \times 10^{-8}$), one honest null (REJECT where no failure existed to improve on), and a self-stop on unmapped state (§4.5).

All results were produced on a single RTX 5090; every number in this paper reproduces from a script in the accompanying repository, and every embodied verdict was pre-registered (decision rule fixed before the runs).

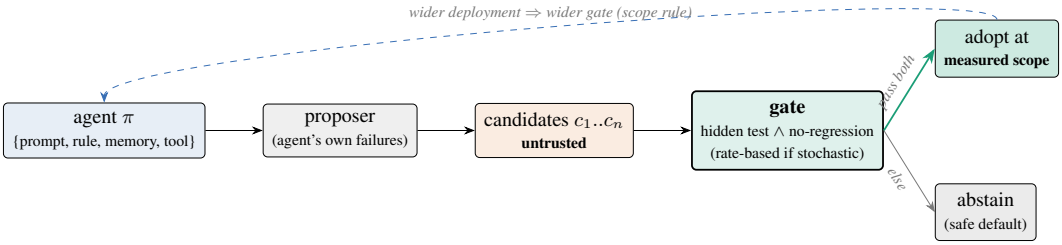

Figure 1: The self-verifying agent. Every self-modification is untrusted until it clears a hidden (held-out, executed where possible) test *and* a regression check over the existing capability profile; otherwise the agent abstains. A verdict is valid only at the scope it was measured—global adoption requires a deployment-scope gate.

2 Method: the self-verifying agent

Let an agent have capabilities measured by evaluators $E = \{e_1, \ldots, e_k\}$ (its *profile*) and let a proposer (the agent itself, prompted over its own failures) emit candidate modifications $c_1, \ldots, c_n$ for a target capability t. The gate (Figure 1):

- **Hidden test:** $\Delta_t(c) = e_t^{\text{held-out}}(\pi + c) - e_t^{\text{held-out}}(\pi)$ must exceed a gain margin, where π is the current agent. The held-out set must be data the candidate was not derived from; on code, *execution* of the candidate on held-out tests is the evaluator.
- **Regression check:** $\min_j [\, e_j(\pi + c) - e_j(\pi) \,] \geq -\varepsilon$ over all existing capabilities j, for noise margin ε.
- **Selection:** among candidates passing both, adopt the best by hidden score; else adopt nothing (abstain).
- **Stochastic domains:** replace scores with failure rates over N runs and require a one-sided significant reduction (exact binomial test). Single runs are meaningless when the baseline itself is flaky (§4.1).
- **Scope rule:** a verdict licenses adoption only at the scope of the evaluators that produced it. Deployment-scope adoption requires a deployment-scope gate: a screen over the deployable distribution, followed by interleaved candidate/baseline confirmation runs on flagged cases under a pre-registered decision rule (§4.3).

Two design lessons we verified empirically. First, the hidden set must be large or refreshed across selection rounds: reusing a 20-item set over 8 candidate selections re-inflated scores by $+0.045$ (optimism from gate reuse)—the system-level analogue of adaptive overfitting to a test set. Second, the proposer is itself unreliable: models rationalize their failures with their priors instead of fitting the evidence, and the induced rules are wrong until the proposer is forced to fit input–output patterns; even then, *verification-as-selection*, not proposer quality, is what makes the loop converge. Both motivate the central design stance: assume the proposal stream is adversarial.

3 Fake improvement is the norm

One pattern repeats across every surface an agent can modify: the score on the data a modification was derived from is uninformative about its value, because the proposal stream itself manufactures candidates that ace exactly that data. We measure this at four levels—fine-tunes, silent regressions, tool rewrites, and real code—then show the two-gate verifier converting each into a clean adoption decision (Figure 2).

Setup. Text experiments use a 7B instruct model (Qwen2.5-7B-Instruct) as both proposer and substrate; the judge/abstention substrate underneath was validated separately (a zero-shot reading judge at 0.949 vs. a tuned guardrail's 0.789 on a sealed benchmark, with calibrated abstention transferring to 4/6 unseen domains). Code experiments execute real Python against held-out unit tests.

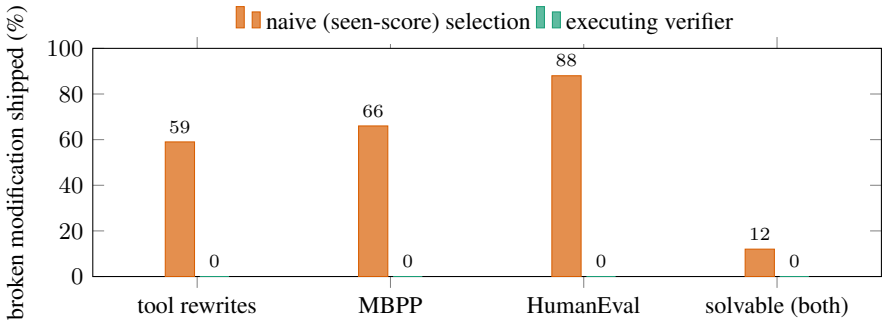

Figure 2: The single adoption decision. When a fake modification exists, naive selection by the score on the data the modification was derived from ships it 59–88% of the time; the verifier (execution on held-out data + regression check) ships 0% across all four settings, while still adopting real improvements (88% of tool decisions; 68% of MBPP problems).

Fine-tuning level. Among three candidate LoRA fine-tunes of the judge, the one trained on a single source memorizes it (train loss 0.000, in-distribution 0.993) and generalizes worst (held-out 0.507). Naive selection (best in-distribution) deploys exactly this memorizer; held-out selection deploys the broad candidate (0.578).

Regression is silent. Fine-tuning to fix one weak source costs another source -0.16 while the target barely moves ($+0.005$); a 1-D "did the target improve?" check misses it entirely. A retention-replay candidate removes the regression (worst old-capability $\Delta = 0.000$). This reproduces, on a judge, the classic naive-correction backfire—and previews the embodied result of §4.3 in miniature.

Tool rewrites. Asked to extend its own Python tool, the agent's proposal set contains a hardcoded lookup of the observed examples in **100% of 24 decisions**; the seen score cannot distinguish it from the general tool (both 1.00). Naive selection therefore ships a broken tool **59%** of the time. *Executing* candidates on held-out inputs ships **0%**, while still adopting a working tool in 88% of decisions.

Real code. On MBPP (n=100) and HumanEval (n=80), a genuine overfit—a solution passing the shown test and failing held-out tests—appears on **29% / 30%** of problems respectively. Where it exists, naive ships it **19/29** and **21/24** times; the verifier **0 on both**. Restricted to problems the agent can solve (removing the abstain-on-unsolvable confound), naive ships a held-out-failing solution **12% on both benchmarks**, the executing verifier **0%**. Two independent benchmarks, near-identical rates: the effect is a property of the self-modification stream, not of a dataset (Figure 2).

3.1 The flagship decision, and the autonomous loop

An agent with capability A (1.00) and target B (0.00) proposes rule-edits from its own failures. The verifier adopts the edit that fixes B *and* keeps A—**rejecting the edit that fixes B best (hidden 1.00) because it regresses A**—ending at A $= 1.00$, B $= 0.83$ (mean 0.92) vs. naive 0.67. The regression check is not decoration: the best target-fix can be the worst adoption. The loop also closes autonomously: forced to fit patterns rather than rationalize, the proposer induces the true hidden rule among its candidates and the verifier selects it (2/2 tasks).

3.2 Honest negative: cumulative compounding

Over 3 acquisition rounds the gated agent leads at every round (0.88 vs. 0.79 final). At 5 rounds the cumulative-mean advantage shrinks to $+0.03$: worst-case capability becomes dominated by new-task acquisition noise, not by regression. The robust claim in this paper is the *per-decision* advantage; large multi-round compounding is not demonstrated on these substrates and we do not claim it.

4 Embodied: the full arc on a real driving system

We instantiate the same gate on CARLA/Bench2Drive [Dosovitskiy et al., 2017, Jia et al., 2024] with a real end-to-end student (the LEAD tfv6 transfuser-family stack) and its privileged expert as teacher. This section is one continuous, pre-registered narrative: flakiness forces a statistical gate (§4.1); the gate rejects a harmful fix and accepts a real one (§4.2); the accepted candidate is then *globally* rejected at deployment scope (§4.3); and a ten-recipe repair search under the fixed gate ends in a mapped frontier and an honest negative (§4.4).

4.1 Failures are flaky; verification must be rate-based

Re-running routes with the agent, weather, and configuration unchanged: route 11755 (an `EnterActorFlow` scenario) fails $7/12 = 58\%$ of runs (95% CI $[0.32, 0.81]$)—the same route flips PASS/FAIL run to run—while three routes labeled FAILURE by the single-run benchmark pass 100% on re-run: their recorded failures were one-off noise. Two consequences. (i) Single-run failure labels, the leaderboard standard, measure stochasticity as much as capability (run-to-run variance on these benchmarks is documented [Jia et al., 2024]; we operationalize it into the gate rather than re-discover it). (ii) Verifying a fix means comparing failure *rates*: confirming a $58\% \to 20\%$ improvement needs $\sim$19 runs per arm at 80% power. A naive single-run check of an unchanged model would "confirm a fix" with 42% probability.

The first live fix candidate makes the point concrete: a plausible control-layer intervention ("drive more conservatively") *raised* the observed failure rate ($50\% \to 83\%$; collisions persisted and slow driving added penalties), and the statistical gate rejected it before deployment.

4.2 The accepted fix: the loop closes

The full cycle then ran end-to-end: student fails 11755 at 58% $\to$ the privileged expert is stable-perfect there $(8/8)$ $\to$ 359 MB of expert demonstrations collected on the failure scenario plus retention routes $\to$ 10-epoch fine-tune from the released checkpoint (single RTX 5090, $\sim$2.5 h) $\to$ the pre-registered gate on fresh rollouts: **hidden test** $0/19$ failures (baseline $7/12$; exact binomial $p = 6.0 \times 10^{-8}$); **regression** $0/12$ on the retention routes $\to$ **ACCEPT**. To our knowledge this is the first pre-registered, error-controlled statistical *accept* produced inside a closed self-improvement loop on an embodied system—and the same harness had already produced the corresponding *reject* (§4.1).

The honest scope of that verdict, stated at the time: verified *same-route repair* over fresh stochastic rollouts plus a 3-route regression set—not cross-route generalization, and not a full regression sweep. That scope statement is not a caveat; it is the next experiment.

4.3 The deployment gate: same candidate, opposite verdict

Is the locally-accepted candidate safe to *deploy*—to swap in globally? The deployment-scope gate answers with a screen over the 169 baseline-clean routes, then interleaved candidate/baseline confirmation runs ($4{+}4$ per flagged route) under a pre-registered rule (confirmed regression: candidate $\geq 3/4$ failures and baseline $\leq 1/4$; global reject if any confirmed regression or pooled candidate failure significantly above baseline).

The screen flagged 19 collision routes plus a systematic soft signature (14 routes at exactly score 70 with a red-light violation and mass minimum-speed infractions—a behavioral warp, not noise). The confirmation batch (176/176 valid interleaved runs) returned: **12 confirmed regressions**; pooled on flagged routes, candidate $53/76 = 70\%$ vs. baseline $5/76 = 7\%$ failures, one-sided $p = 8.5 \times 10^{-45}$ $\to$ **GLOBAL REJECT** (Figure 3).

The mechanism is classic catastrophic forgetting [McCloskey and Cohen, 1989, French, 1999]: ten epochs on $\sim$1.6k frames of four scenario types with a 3-route retention set warped global driving—collisions on ordinary routes, red-light violations, slow driving everywhere. What is new here is not the phenomenon but the demonstration: *the same candidate earns a legitimate statistical ACCEPT at target scope and a $p = 8.5 \times 10^{-45}$ REJECT at deployment scope, on a real system, under pre-registered rules*. Single-target evaluation—the standard practice when a fix is verified at all—would have shipped this model. The 3-route retention check passed $(0/12)$ because those routes happen to sit among the least-affected; a narrow regression set does not just under-measure damage, it can

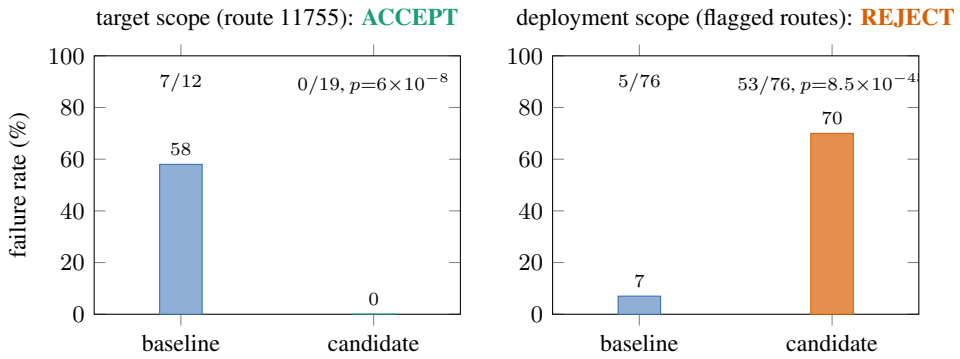

Figure 3: One candidate, two scopes, opposite verdicts—both correct at their scope. Left: the fine-tune genuinely repairs its target route ($58\% \to 0\%$ over 19 fresh rollouts). Right: the same model at deployment scope collides on ordinary routes the baseline drives cleanly (12 confirmed regressions; pooled 70% vs. 7%). Single-target evaluation would have shipped it.

Table 1: The repair-search ledger. Every verdict pre-registered; the gate never adopted a regression. Fix line: $\leq 2/8$ failures on route 11755. Retention line: $\leq 30\%$ pooled failures on the 12 confirmed-regression routes. †K1's retention was measured at deployment scope (§4.3); panel lines were pre-registered after it.

#	recipe (one variable at a time where possible)	fix 11755	retention	verdict
0	control-layer "drive conservatively" (no training)	worse ($50 \to 83\%$)	—	REJECT
K1	narrow data, 10 epochs, full FT	$0/19$ ✓	12 confirmed reg.†	**global REJECT**
K2a	K1 data, 3 epochs	$0/8$ ✓	$12/12$ routes fail	REJECT
K2b	broad retention (17 routes), 3 ep	$4/8$	33%	REJECT (fix diluted)
A1	fix 40% oversample, 3 ep	$4/8$	25% ✓	REJECT (fix diluted)
A2	fix 60% oversample, 5 ep	$0/8$ ✓	29% ✓	**panel PASS →**
	↪ deployment gate	—	6 confirmed reg., $p=2.5\times10^{-20}$	**global REJECT**
A3	A2 + fix-family demos in fix bucket	$6/7$	(cut short)	REJECT (family conflict)
A4	A2 fix bucket + targeted retention	$4/8$	33%	REJECT
A5	A4 + frozen backbone (head-only)	$4/8$	$\sim 31\%$	REJECT
A6	A2 data + frozen backbone	$3/8$	$\mathbf{17}\%$ ✓	REJECT (fix, by one run)

systematically miss it. **Verification scope must equal deployment scope.** Nothing was adopted; the deployed capability floor never dropped.

4.4 The ten-recipe repair search: a mapped frontier, an honest negative, a gate that held

Can a better fine-tuning recipe pass *both* gates? We ran a pre-registered, timeboxed search: the gate stays fixed; only the recipe iterates. Local panel (pre-registered lines): fix arm 11755 $\leq 2/8$ failures; retention arm over the 12 previously-confirmed regression routes, pooled $\leq 30\%$; deployment gate as in §4.3 for any panel-passer. Ten recipes were evaluated in total (Table 1; $\sim$450 verification rollouts): the control-layer intervention, the original narrow fine-tune, and eight variants spanning training length, retention breadth, mixture ratios (via the stack's native two-bucket oversampling), fix-family breadth, targeted retention collection, and backbone freezing (head-only, 16.6M trainable parameters).

The search converged, honestly, to a mapped trade-off rather than a winner (Figure 4). Its findings:

- **Forgetting here is a data-breadth problem, not a training-length problem**: cutting epochs $10 \to 3$ on narrow data preserved the repair and all twelve regressions.
- **The stability–plasticity see-saw is real and sits in the planning head**: with fix data identical to the panel-passing recipe, merely *adding* retention data broke the repair (A4); freezing the perception backbone did not recover it (A5); the interference lives in the shared planning parameters, echoing at system scale the classic trade-off [French, 1999].

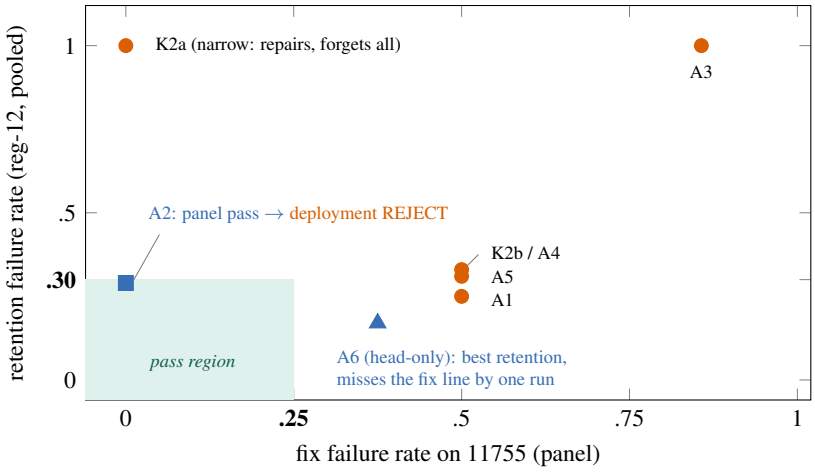

Figure 4: The measured repair–retention frontier across the recipe search (panel scope; A3's retention arm was cut short and is plotted at its decided axis only). Narrow full fine-tuning repairs and forgets (upper left); breadth or head-only training retains and dilutes (lower right). A2 enters the pass region at panel scope and is then rejected by the deployment gate—the scope rule again. No recipe passes both gates at this budget.

- **Dilution is fractal**: adding same-family demonstrations to the fix bucket did not generalize the repair—it diluted and conflicted with it one level down (A3: the family-scenario route itself had failed $4/4$ under A2, a within-family overfit; fixing that broke 11755).
- **Panel pass $\neq$ deployment pass, again**: A2, the one panel-passer, was rejected by the deployment gate (6 confirmed regressions, $p = 2.5 \times 10^{-20}$)—though the search was converging (12 confirmed regressions $\rightarrow$ 6, global behavioral warp eliminated).
- **Bounded negative**: at this budget (a single consumer GPU, the stack's native fine-tuning machinery, $\sim$450 verification rollouts), no recipe passed both gates. Full fine-tuning repairs perfectly and collapses globally; head-only training retains best (17%) and misses the fix line by one run ($3/8$). We report this as a first-class result, not a failure of the method under test: the method under test is the *gate*, and the gate did exactly its job.

Through the entire search—ten recipes, every verdict pre-registered—**zero regressions were adopted and the deployed capability floor never dropped**. The gate demonstrated all four of its decision modes on a real system: accept-real (§4.2), reject-harmful (§4.1, §4.3), reject-insufficient (A6), and protect-through-search (this section). A pre-registered caution for practitioners: iterating recipes against a fixed screen risks overfitting the gate itself (the embodied form of the gate-reuse effect of §2); any finally-accepted candidate must additionally pass held-back validation never used during iteration.

4.5 The outer loop, automated

The searches above were orchestrated by a human (verdicts were always computed by the gate). A fair objection: is this a self-improvement *loop*, or a human-curated search? We answer by removing the human. Under a pre-registered protocol (win conditions, a hard cutoff, and a compute fallback, all sealed before any code was written), we implemented the outer loop as code: an *append-only trial registry* (thresholds stored at registration time, before data; trials frozen once a verdict is computed; verdicts computed from raw counts, never stored by fiat) and a *rule-based prescriber* (the continual-learning ladder: cheapest variable $\rightarrow$ replay breadth $\rightarrow$ replay weighting $\rightarrow$ targeted replay $\rightarrow$ parameter isolation) that may only propose—the gate remains the sole adoption authority.

Replay (consistency). Replaying the S1 search through this machinery reproduces all 12 recorded verdicts exactly from raw counts—including the scope splits and the honest early-stops—and the prescriber emits the move the human search actually made at 9/9 recipe transitions, both scope

escalations, and the budget-exhausted closure. Sealed caveat: the rules were distilled *from* that search, so replay validates the distillation, not generality.

New trials (zero human verdicts). The automated loop then ran fresh trials end-to-end (the human launched a script and read the result). On 100 never-before-used MBPP problems, with each problem's tests split visible / gate-hidden / *audit* (the gate never sees the audit; without this split "the gate ships zero failures" would be true by construction): naive ships an audit-failing solution $41/100$; the gate adopts 59, abstains 41, and its adoptions fail audit $5/59 = 8.5\%$—a registry-computed **ACCEPT** ($p = 2.8 \times 10^{-8}$). Note the honest non-zero: the audit measures what the gate buys ($41\% \to 8.5\%$), not a tautology. The prescriber—not a human—then ordered scope escalation, and on the remaining 26 fresh HumanEval problems naive shipped *zero* audit failures: the registry returned **REJECT** (no improvement evidence, $p = 1$), declining to manufacture significance where none exists, and the prescriber, hitting a state its rules do not map, handed back to the operator and stopped rather than guessing. One accept, one honest null, no human in the verdict path: the title's claim is literal for these trials.

5 What is and is not established

Established. (1) The agent's own modification stream is adversarial by default—overfits and lookup-hacks are pervasive across surfaces and domains, including real code. (2) Seen/derivation-set score carries no signal against them; held-out (on code, *executed*) evaluation does. (3) The two-gate verifier converts this into a reliable per-decision win: $59\% \to 0\%$, $19/29 \to 0$, $21/24 \to 0$. (4) Regression checking is necessary: the best target-fix can be the worst adoption. (5) In stochastic embodied domains, only rate-based statistical verification is meaningful, and it supports both accept ($p = 6 \times 10^{-8}$) and reject ($p = 8.5 \times 10^{-45}$) with error control. (6) A verdict is only valid at its measured scope; scope-mismatched acceptance ships catastrophic forgetting on a real system. (7) The outer loop needs no human: registry + prescriber + gate replays the recorded search exactly from raw counts and runs fresh trials end-to-end, accepting where evidence exists and returning an honest null where it does not (§4.5).

Not established. Large cumulative multi-round gains (§3.2); verified self-improvement on hard tasks where no candidate transfers (the verifier then correctly adopts nothing—safe, but not improvement); a deployment-gate-passing embodied repair (§4.4: bounded negative at this budget); behavior beyond a 7B proposer, benchmark-scale substrates, one vehicle stack, and one simulator.

6 Related work and claim boundaries

Self-improving agents (Reflexion [Shinn et al., 2023]; ExpeL [Zhao et al., 2024]; AutoGuide [Fu et al., 2024]; Voyager [Wang et al., 2024]) supply the proposer we gate. **That self-generated improvement is frequently fake is established**: agent benchmarks are reward-hackable [Wang et al., 2026]; frontier agents hack their evaluations spontaneously [METR, 2025]; a self-modifying system falsified its own metric [Zhang et al., 2025]; self-reflection confabulates. **Gating self-modification is an emerging thread**: PACE shows that greedy "keep it if the score went up" acceptance is uncontrolled adaptive multiple testing and replaces it with anytime-valid (e-process) commit tests on text-agent loops [Shawn, 2026]—the same diagnosis as ours, in the statistical dimension we pre-register; GRASP gates skill-library edits on probe-based net improvement [GRASP authors, 2026]; GRACE uses an update-rejection gate inside prompt optimization [GRACE authors, 2025]; and sequential statistics have been brought to robot policy comparison as an offline evaluation tool [Snyder et al., 2026]. None of these operate on an embodied system inside the loop, and none demonstrate the deployment-scope failure. **Driving-benchmark noise is documented** [Jia et al., 2024, Dosovitskiy et al., 2017], with published scores typically single-run. **Continual learning** names our regression phenomenon catastrophic forgetting [McCloskey and Cohen, 1989, French, 1999, Kirkpatrick et al., 2017], and our repair recipes (replay ratios, freezing, DAgger-style collection [Ross et al., 2011, Hu et al., 2022]) are its textbook tools deliberately reused, not claimed as new.

We therefore *do not* claim: the first observation that self-improvement is often fake; the first statistical gate for self-improving agents; or the discovery of driving-benchmark flakiness. To our knowledge, what *is* unclaimed before this work: (a) fake-improvement *rates measured across all four modification surfaces* under one protocol, plus two real code benchmarks; (b) a *pre-registered, error-controlled*

statistical accept and reject operating inside a closed self-improvement loop on an embodied system, including the operationalization of run-to-run noise into a failure-rate gate; and (c) the *deployment-scope demonstration*—a narrowly-scoped gate accepts a change whose collateral damage a scope-matched gate then catches, on a real system, followed by a gate-protected recipe search ending in a mapped frontier. (Literature scan of July 2026; the space moves fast.)

7 Limitations

A 7B proposer; controlled/benchmark substrates; six-item held-out sets on the toy loops (granularity 0.17); one vehicle stack and one simulator; the embodied fix targets one (representative, flakiest) failure route with an 8-run fix arm and 24-run retention arm at panel scope—adequate for the pre-registered lines, underpowered for small effects; the bounded negative of §4.4 is a statement about this training budget and recipe family (parameter-efficient adapters and larger fix-side data remain untested). Gate reuse must be budgeted (§2); we spec but do not automate it. Multi-round compounding may appear on richer substrates; we did not demonstrate it.

8 Conclusion

Across text, real code, and embodied driving, self-improvement without verification ships fakes at rates of 12–88%, and the fakes are supplied by the loop itself. A verifier that demands held-out transfer and no regression—nothing more exotic—reduced shipped fakes to zero in our experiments while still adopting real improvements; its statistical form both accepted a real embodied repair and rejected the same candidate at deployment scope before it could ship catastrophic forgetting; and under a ten-recipe search it held the capability floor while honestly reporting that no recipe earned adoption. If self-improving agents are to be trusted with their own modification stream, the verifier is not a component of the loop. **The verifier *is* the loop; the rest is a proposal generator.**

References

Alexey Dosovitskiy, German Ros, Felipe Codevilla, Antonio Lopez, and Vladlen Koltun. CARLA: An open urban driving simulator. In *Proceedings of the 1st Annual Conference on Robot Learning (CoRL)*, 2017.

Robert M. French. Catastrophic forgetting in connectionist networks. *Trends in Cognitive Sciences*, 3 (4):128–135, 1999.

Yao Fu, Dong-Ki Kim, Jaekyeom Kim, Sungryull Sohn, Lajanugen Logeswaran, Kyunghoon Bae, and Honglak Lee. Autoguide: Automated generation and selection of context-aware guidelines for large language model agents. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.

GRACE authors. GRACE: Gated refinement and adaptive compression for prompt optimization. *arXiv preprint arXiv:2509.23387*, 2025. Re-verify exact title and author list at camera-ready.

GRASP authors. GRASP: Gated regression-aware skill proposer for self-improving LLM agents. *arXiv preprint arXiv:2605.29668*, 2026. Re-verify author list at camera-ready.

Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations (ICLR)*, 2022.

Xiaosong Jia, Zhenjie Yang, Qifeng Li, Zhiyuan Zhang, and Junchi Yan. Bench2drive: Towards multi-ability benchmarking of closed-loop end-to-end autonomous driving. In *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks Track*, 2024.

James Kirkpatrick, Razvan Pascanu, Neil Rabinowitz, Joel Veness, Guillaume Desjardins, Andrei A. Rusu, Kieran Milan, John Quan, Tiago Ramalho, Agnieszka Grabska-Barwinska, Demis Hassabis, Claudia Clopath, Dharshan Kumaran, and Raia Hadsell. Overcoming catastrophic forgetting in neural networks. *Proceedings of the National Academy of Sciences*, 114(13):3521–3526, 2017.

Michael McCloskey and Neal J. Cohen. Catastrophic interference in connectionist networks: The sequential learning problem. *Psychology of Learning and Motivation*, 24:109–165, 1989.

METR. Recent frontier models are reward hacking. `https://metr.org/blog/2025-06-05-recent-reward-hacking/`, 2025.

Stéphane Ross, Geoffrey Gordon, and Drew Bagnell. A reduction of imitation learning and structured prediction to no-regret online learning. In *Proceedings of the 14th International Conference on Artificial Intelligence and Statistics (AISTATS)*, 2011.

Zayx Shawn. PACE: Anytime-valid acceptance tests for self-evolving agents. *arXiv preprint arXiv:2606.08106*, 2026. Author spelling per arXiv listing; re-verify at camera-ready.

Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: Language agents with verbal reinforcement learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.

David Snyder, Apurva Badithela, Nikolai Matni, George Pappas, Anirudha Majumdar, Masha Itkina, and Haruki Nishimura. Beyond binary success: Sample-efficient and statistically rigorous robot policy comparison. *arXiv preprint arXiv:2603.13616*, 2026.

Guanzhi Wang, Yuqi Xie, Yunfan Jiang, Ajay Mandlekar, Chaowei Xiao, Yuke Zhu, Linxi Fan, and Anima Anandkumar. Voyager: An open-ended embodied agent with large language models. *Transactions on Machine Learning Research*, 2024.

Hao Wang, Qiuyang Mang, Alvin Cheung, Koushik Sen, and Dawn Song. Do androids dream of breaking the game? systematically auditing AI agent benchmarks with BenchJack. *arXiv preprint arXiv:2605.12673*, 2026.

Jenny Zhang, Shengran Hu, Cong Lu, Robert Lange, and Jeff Clune. Darwin gödel machine: Open-ended evolution of self-improving agents. *arXiv preprint arXiv:2505.22954*, 2025.

Andrew Zhao, Daniel Huang, Quentin Xu, Matthieu Lin, Yong-Jin Liu, and Gao Huang. Expel: Llm agents are experiential learners. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2024.

A Reproducibility

All numbers in this paper are produced by scripts in the accompanying repository (anonymized for review); each result entry in the repository's ledger names the script and exact configuration that produced it. Text and code experiments run on a single GPU with a 7B open-weight model; embodied experiments use CARLA 0.9.15 + Bench2Drive with the LEAD stack on a single RTX 5090 (Windows-native). The gate itself is released as a dependency-free library (`vsi_gate.py`) exposing `gate()` (score form) and `gate_rate()` (exact-binomial failure-rate form), with a deterministic no-LLM demo.

B Pre-registered decision rules (embodied)

Fixed before the corresponding runs, unchanged throughout: **(local panel)** fix arm: route 11755, 8 fresh rollouts, pass iff ≤ 2 failures; retention arm: the 12 confirmed-regression routes $\times 2$, pass iff pooled failures $\leq 30\%$; both required. **(deployment gate)** screen all 169 baseline-clean routes once; flag any failure; triage pure-artifact flags (minimum-speed/lane-only) out; confirm each flagged route with interleaved candidate/baseline runs ($4{+}4$); confirmed regression iff candidate $\geq 3/4$ failures and baseline $\leq 1/4$; global REJECT iff ≥ 1 confirmed regression or pooled candidate failure significantly above baseline (one-sided exact test, $\alpha = 0.05$). **(rate gate)** a fix claim requires a one-sided exact-binomial significant reduction against the measured baseline rate, power-planned before collection (e.g. $58\% \to 20\%$ requires $\sim$19 runs/arm at 80% power).